

CS221 Fall 2018 Homework 4

SUNet ID: 05794739

Name: Luis Perez

Collaborators:

By turning in this assignment, I agree by the Stanford honor code and declare that all of this is my own work.

Problem 1

- (a) We give the value for each iteration. We note that $V_{\text{opt}}^0(s) = 0$ to start out. We also note that since for $s_t \in \{-2, 2\}$ we are at a terminal state, we'll have $V_{\text{opt}}(s_t) = 0$ for all iterations.

- (a) After iteration 0, we'll have:

$$V_{\text{opt}}^0(-1) = 0$$

$$V_{\text{opt}}^0(0) = 0$$

$$V_{\text{opt}}^0(1) = 0$$

- (b) After the first iteration, we'll have the following values:

$$\begin{aligned} V_{\text{opt}}^1(-1) &= \max_{a \in \{-1, 1\}} \{0.8[20 + V_{\text{opt}}^0(-2)] + 0.2[-5 + V_{\text{opt}}^0(0)], 0.7[20 + V_{\text{opt}}^0(-2)] + 0.3[-5 + V_{\text{opt}}^0(0)]\} \\ &= 15 \end{aligned}$$

$$\begin{aligned} V_{\text{opt}}^1(0) &= \max_{a \in \{-1, 1\}} \{0.8[-5 + V_{\text{opt}}^0(-1)] + 0.2[-5 + V_{\text{opt}}^0(1)], 0.7[-5 + V_{\text{opt}}^0(-1)] + 0.3[-5 + V_{\text{opt}}^0(1)]\} \\ &= -5 \end{aligned}$$

$$\begin{aligned} V_{\text{opt}}^1(1) &= \max_{a \in \{-1, 1\}} \{0.8[-5 + V_{\text{opt}}^0(0)] + 0.2[100 + V_{\text{opt}}^0(2)], 0.7[-5 + V_{\text{opt}}^0(0)] + 0.3[100 + V_{\text{opt}}^0(2)]\} \\ &= 26.5 \end{aligned}$$

- (c) Finally, after the second iteration, we'll have:

$$\begin{aligned} V_{\text{opt}}^2(-1) &= \max_{a \in \{-1, 1\}} \{0.8[20 + V_{\text{opt}}^1(-2)] + 0.2[-5 + V_{\text{opt}}^1(0)], 0.7[20 + V_{\text{opt}}^1(-2)] + 0.3[-5 + V_{\text{opt}}^1(0)]\} \\ &= 14 \end{aligned}$$

$$\begin{aligned} V_{\text{opt}}^2(0) &= \max_{a \in \{-1, 1\}} \{0.8[-5 + V_{\text{opt}}^1(-1)] + 0.2[-5 + V_{\text{opt}}^1(1)], 0.7[-5 + V_{\text{opt}}^1(-1)] + 0.3[-5 + V_{\text{opt}}^1(1)]\} \\ &= 13.45 \end{aligned}$$

$$\begin{aligned} V_{\text{opt}}^2(1) &= \max_{a \in \{-1, 1\}} \{0.8[-5 + V_{\text{opt}}^1(0)] + 0.2[100 + V_{\text{opt}}^1(2)], 0.7[-5 + V_{\text{opt}}^1(0)] + 0.3[100 + V_{\text{opt}}^1(2)]\} \\ &= 23 \end{aligned}$$

- (b) We interpret this question as asking for the resulting optimal policy for non-terminal states after two iterations. In that case, we have:

$$\begin{aligned}
\pi_{\text{opt}}^2(-1) &= \arg \max_{a \in \{-1,1\}} \{0.8[20 + V_{\text{opt}}^1(-2)] + 0.2[-5 + V_{\text{opt}}^1(0)], 0.7[20 + V_{\text{opt}}^1(-2)] + 0.3[-5 + V_{\text{opt}}^1(0)]\} \\
&= -1 \\
\pi_{\text{opt}}^2(0) &= \arg \max_{a \in \{-1,1\}} \{0.8[-5 + V_{\text{opt}}^1(-1)] + 0.2[-5 + V_{\text{opt}}^1(1)], 0.7[-5 + V_{\text{opt}}^1(-1)] + 0.3[-5 + V_{\text{opt}}^1(1)]\} \\
&= 1 \\
\pi_{\text{opt}}^2(1) &= \arg \max_{a \in \{-1,1\}} \{0.8[-5 + V_{\text{opt}}^1(0)] + 0.2[100 + V_{\text{opt}}^1(2)], 0.7[-5 + V_{\text{opt}}^1(0)] + 0.3[100 + V_{\text{opt}}^1(2)]\} \\
&= 1
\end{aligned}$$

Problem 2

- (a) It is not always the case that $V_1(s_{\text{start}}) \geq V_2(s_{\text{start}})$. For a counter-examples, see “submission.py”.
- (b) The algorithm is rather straight-forward in the case where we have an acyclic MDP.
- The first-step in the algorithm is to topologically sort the graph. It is well-known that for a DAG, a topological sorting is possible and can be computed by a modified version of DFS in linear time ¹
 - Once we have this topological sorting of the states, we process each state in reverse-topological order and compute

$$V(s) = \max_{a \in A} \left\{ \sum_{s' \in \text{Succ}(s,a)} T(s, a, s') [R(s, a, s') + V(s')] \right\}$$

directly for each such state.

- After this single pass over the states, we return the resulting value function.

We claim that the computed $V(s) = V_{\text{opt}}(s)$ (ie, in this single pass, we’ve computed the optimal value function). To understand why, we must recall that a topological sorting is one such that for every edge transition (s, a, s') , from s to s' , s comes before s' in the ordering. In our algorithm above, we processed these states in reverse order (ie, we calculate the value of s' before we compute the value of s). More formally, consider all terminal states (ie, states with no successors). These states are processed first by our algorithm, given us the base case:

$$V(s') = 0 = V_{\text{opt}}(s') \quad (\text{for all terminal states } s')$$

¹https://en.wikipedia.org/wiki/Topological_sorting#Depth-first_search

Now let us assume $V(s') = V_{\text{opt}}(s')$ for all s' which our algorithm has already processed (ie, if our algorithm is processing states s , then the above holds true for all states s' which fall after s in the topological sort). Consider the processing of state s . For this state, our algorithm will compute:

$$V(s) = \max_{a \in A} \left\{ \sum_{s' \in \text{Succ}(s,a)} T(s, a, s') [R(s, a, s') + V(s')] \right\} \quad (\text{definition of our algorithm})$$

$$= \max_{a \in A} \left\{ \sum_{s' \in \text{Succ}(s,a)} T(s, a, s') [R(s, a, s') + V_{\text{opt}}(s')] \right\}$$

(all s' are descendants of s , and therefore, by the inductive hypothesis, we have $V(s') = V_{\text{opt}}(s')$)
 $= V_{\text{opt}}(s)$ (definition of V_{opt})

- (c) Following the hint, the solution to this problem is essentially given to us in lecture. The problem already provides States', Actions'(s), and γ' . As per Percy's lecture notes, we define the transition probabilities and rewards as follows:

$$T'(s, a, s') = \begin{cases} (1 - \gamma) & s' = o \\ \gamma T(s, a, s') & \text{otherwise} \end{cases}$$

$$R'(s, a, s') = \begin{cases} 0 & s' = o \\ R(s, a, s') & \text{otherwise} \end{cases}$$

Informally, with probability $(1 - \gamma)$ every state can now end in a terminal state with (receiving 0 reward). All other transitions probabilities are discounted by γ . We claim that $V_{\text{opt}}(s) = V'_{\text{opt}}$ for all $s \in \text{States}$. We can prove this directly. First, let's recall that if $V_{\text{opt}}(s)$ exists, it is the unique solution to:

$$V_{\text{opt}}(s) = \max_{a \in \text{Actions}(s)} \left\{ \sum_{s' \in \text{Succ}(s,a)} T(s, a, s') [R(s, a, s') + V_{\text{opt}}(s')] \right\}$$

Now let us consider a state $s \in \text{States}$. Then we have:

$$\begin{aligned}
V'_{\text{opt}}(s) &= \max_{a \in \text{Actions}'(s)} \left\{ \sum_{s' \in \text{Succ}'(s,a)} T'(s, a, s') [R(s, a, s') + V'_{\text{opt}}(s')] \right\} \quad (\text{definition of } V'_{\text{opt}}) \\
&= \max_{a \in \text{Actions}(s)} \left\{ T'(s, a, o) [R'(s, a, o) + V'_{\text{opt}}(o)] + \sum_{s' \in \text{Succ}(s,a)} T'(s, a, s') [R'(s, a, s') + V'_{\text{opt}}(s')] \right\} \\
&\quad (\text{Actions}'(s) = \text{Actions}(s) \text{ and } \text{Succ}'(s, a) = \{o\} \cup \text{Succ}(s, a) \text{ by contruction}) \\
&= \max_{a \in \text{Actions}(s)} \left\{ \sum_{s' \in \text{Succ}(s,a)} \gamma T(s, a, s') [R(s, a, s') + V'_{\text{opt}}(s')] \right\}. \\
&\quad (R'(s, a, o) + V'_{\text{opt}}(o) = 0 \text{ and definition of } R' \text{ and } T')
\end{aligned}$$

Note that the above equation is precisely the equation that $V_{\text{opt}}(s)$ solves. Therefore, we must have that $V_{\text{opt}}(s)$ and $V'_{\text{opt}}(s)$ are the same function.